

## Phase 1: Dependencies and Setup

```
# Install all required dependencies
!pip install torch torchvision torchaudio
!pip install opencv-python-headless
!pip install pillow
!pip install scikit-image
!pip install tqdm
!pip install ipywidgets

# Fix for PyTorch 2.6+ compatibility
!pip install huggingface_hub==0.25.2

# Import all required libraries
import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
from torch.utils.data import Dataset, DataLoader
import torchvision.transforms as transforms
import cv2
import numpy as np
from PIL import Image
import os
import zipfile
import urllib.request
from tqdm import tqdm
import matplotlib.pyplot as plt
from skimage.metrics import structural_similarity as ssim
import json
import time
from google.colab import files
import shutil
import glob
import base64
from IPython.display import display, HTML
import ipywidgets as widgets

# Set device
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"Using device: {device}")
print(f"CUDA available: {torch.cuda.is_available()}")
if torch.cuda.is_available():
    print(f"GPU: {torch.cuda.get_device_name(0)}")
```



◆ What can I help you build?



```

Requirement already satisfied: idna>=2.8 in /usr/local/lib/python3.11/dist-packages (from anyio>=3.1.0->jupyter)
Requirement already satisfied: sniffio>=1.1 in /usr/local/lib/python3.11/dist-packages (from anyio>=3.1.0->jupyter)
Downloading jedi-0.19.2-py2.py3-none-any.whl (1.6 MB)
1.6/1.6 MB 43.4 MB/s eta 0:00:00
Installing collected packages: jedi
Successfully installed jedi-0.19.2
Collecting huggingface_hub==0.25.2
  Downloading huggingface_hub-0.25.2-py3-none-any.whl.metadata (13 kB)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from huggingface_hub==0.25.2)
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.11/dist-packages (from huggingface_hub==0.25.2)
Requirement already satisfied: packaging>=20.9 in /usr/local/lib/python3.11/dist-packages (from huggingface_hub==0.25.2)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.11/dist-packages (from huggingface_hub==0.25.2)
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from huggingface_hub==0.25.2)
Requirement already satisfied: tqdm>=4.42.1 in /usr/local/lib/python3.11/dist-packages (from huggingface_hub==0.25.2)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.11/dist-packages (from huggingface_hub==0.25.2)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests->huggingface_hub)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests->huggingface_hub)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests->huggingface_hub)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests->huggingface_hub)
Downloading huggingface_hub-0.25.2-py3-none-any.whl (436 kB)
436.6/436.6 kB 18.0 MB/s eta 0:00:00
Installing collected packages: huggingface_hub
  Attempting uninstall: huggingface_hub
    Found existing installation: huggingface-hub 0.33.2
    Uninstalling huggingface-hub-0.33.2:
      Successfully uninstalled huggingface-hub-0.33.2
ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This
transformers 4.53.0 requires huggingface-hub<1.0,>=0.30.0, but you have huggingface-hub 0.25.2 which is incompatible.
diffusers 0.34.0 requires huggingface-hub>=0.27.0, but you have huggingface-hub 0.25.2 which is incompatible.
gradio 5.31.0 requires huggingface-hub>=0.28.1, but you have huggingface-hub 0.25.2 which is incompatible.
Successfully installed huggingface_hub-0.25.2
Using device: cpu
CUDA available: False

```

## Phase 2: Model Architectures

# Student Model Architecture (Lightweight CNN)

```

class StudentModel(nn.Module):
    def __init__(self, scale_factor=4):
        super(StudentModel, self).__init__()
        self.scale_factor = scale_factor

        # Feature extraction layers
        self.conv1 = nn.Conv2d(3, 32, kernel_size=3, padding=1)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, padding=1)
        self.conv3 = nn.Conv2d(64, 64, kernel_size=3, padding=1)

        # Residual blocks
        self.res_blocks = nn.ModuleList([
            self._make_residual_block(64) for _ in range(4)
        ])

        # Upsampling layers
        self.upsample = nn.Sequential(
            nn.Conv2d(64, 64 * (scale_factor ** 2), kernel_size=3, padding=1),
            nn.PixelShuffle(scale_factor),
            nn.Conv2d(64, 3, kernel_size=3, padding=1)
        )

        self.relu = nn.ReLU(inplace=True)
        self.tanh = nn.Tanh()

    def _make_residual_block(self, channels):
        return nn.Sequential(
            nn.Conv2d(channels, channels, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.Conv2d(channels, channels, kernel_size=3, padding=1)
        )

    def forward(self, x):

```

```

# Feature extraction
x1 = self.relu(self.conv1(x))
x2 = self.relu(self.conv2(x1))
x3 = self.relu(self.conv3(x2))

# Residual blocks
residual = x3
for res_block in self.res_blocks:
    out = res_block(residual)
    residual = residual + out

# Upsampling
out = self.upsample(residual)
out = self.tanh(out)

return out

# Teacher Model: SRResNet Architecture
class SRResNetTeacherModel:
    def __init__(self, scale=4):
        self.scale = scale
        self.device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

    # Create SRResNet teacher model
    self.model = self._create_srresnet_model().to(self.device)
    self.model.eval()
    print("SRResNet teacher model initialized")

    def _create_srresnet_model(self):
        """Create SRResNet-like architecture"""
        class ResidualBlock(nn.Module):
            def __init__(self, channels):
                super().__init__()
                self.conv1 = nn.Conv2d(channels, channels, 3, padding=1)
                self.bn1 = nn.BatchNorm2d(channels)
                self.conv2 = nn.Conv2d(channels, channels, 3, padding=1)
                self.bn2 = nn.BatchNorm2d(channels)
                self.relu = nn.ReLU(inplace=True)

            def forward(self, x):
                residual = x
                out = self.relu(self.bn1(self.conv1(x)))
                out = self.bn2(self.conv2(out))
                out += residual
                return out

        class SRResNet(nn.Module):
            def __init__(self, scale_factor=4, num_channels=3, num_features=64, num_blocks=16):
                super().__init__()
                self.scale_factor = scale_factor

                # Initial convolution
                self.conv_input = nn.Conv2d(num_channels, num_features, 9, padding=4)

                # Residual blocks
                self.residual_blocks = nn.Sequential(
                    *[ResidualBlock(num_features) for _ in range(num_blocks)]
                )

                # Post-residual convolution
                self.conv_mid = nn.Conv2d(num_features, num_features, 3, padding=1)
                self.bn_mid = nn.BatchNorm2d(num_features)

                # Upsampling layers
                upsampling_layers = []
                for _ in range(int(np.log2(scale_factor))):
                    upsampling_layers.extend([
                        nn.Conv2d(num_features, num_features * 4, 3, padding=1),
                        nn.PixelShuffle(2),

```

```

        nn.ReLU(inplace=True)
    })
    self.upsampling = nn.Sequential(*upsampling_layers)

    # Final convolution
    self.conv_output = nn.Conv2d(num_features, num_channels, 9, padding=4)

    self.relu = nn.ReLU(inplace=True)

    def forward(self, x):
        out = self.relu(self.conv_input(x))
        residual = out
        out = self.residual_blocks(out)
        out = self.bn_mid(self.conv_mid(out))
        out += residual
        out = self.upsampling(out)
        out = self.conv_output(out)
        return torch.tanh(out)

    return SRResNet(scale_factor=self.scale)

def enhance(self, img):
    """Enhance image using SRResNet"""
    try:
        if isinstance(img, np.ndarray):
            # Convert numpy to tensor
            if len(img.shape) == 3:
                img_tensor = torch.from_numpy(img).float() / 255.0
                img_tensor = img_tensor.permute(2, 0, 1).unsqueeze(0)
            else:
                img_tensor = torch.from_numpy(img).float() / 255.0
                img_tensor = img_tensor.unsqueeze(0).unsqueeze(0)
                img_tensor = img_tensor.repeat(1, 3, 1, 1)
        else:
            img_tensor = img.unsqueeze(0) if len(img.shape) == 3 else img

        img_tensor = img_tensor.to(self.device)

        with torch.no_grad():
            enhanced = self.model(img_tensor)
            enhanced = enhanced.squeeze(0).permute(1, 2, 0).cpu().numpy()
            enhanced = np.clip((enhanced + 1) * 127.5, 0, 255).astype(np.uint8)

        return enhanced

    except Exception as e:
        print(f"SRResNet enhancement failed: {e}")
        # Fallback to bicubic upsampling
        h, w = img.shape[:2]
        return cv2.resize(img, (w*self.scale, h*self.scale), interpolation=cv2.INTER_CUBIC)

# Initialize models
student_model = StudentModel(scale_factor=4).to(device)
teacher_model = SRResNetTeacherModel(scale=4)

print(f"Student model parameters: {sum(p.numel() for p in student_model.parameters())},}")
print("Models initialized successfully!")

```

```

🔄 SRResNet teacher model initialized
Student model parameters: 944,323
Models initialized successfully!

```

### Phase 3: Dataset Preparation

```

# Download and setup DIV2K Dataset
def download_div2k_dataset():
    """Download and extract DIV2K dataset"""

```

```

print("Starting DIV2K dataset download...")

if not os.path.exists('DIV2K'):
    os.makedirs('DIV2K')

# URLs for DIV2K dataset
urls = {
    'train_hr': 'http://data.vision.ee.ethz.ch/cvl/DIV2K/DIV2K_train_HR.zip',
    'valid_hr': 'http://data.vision.ee.ethz.ch/cvl/DIV2K/DIV2K_valid_HR.zip'
}

# Download files
for dataset_type, url in urls.items():
    zip_path = f'DIV2K/DIV2K_{dataset_type.replace("_", "-").zip}'

    if not os.path.exists(zip_path):
        print(f'Downloading {dataset_type} dataset...')
        try:
            urllib.request.urlretrieve(url, zip_path)
            print(f'✓ Downloaded {dataset_type}')
        except Exception as e:
            print(f'x Failed to download {dataset_type}: {e}')
            continue
    else:
        print(f'✓ {dataset_type} already exists')

# Extract files
for dataset_type in ['train_hr', 'valid_hr']:
    zip_path = f'DIV2K/DIV2K_{dataset_type.replace("_", "-").zip}'
    extract_path = f'DIV2K/DIV2K_{dataset_type.replace("_", "-")}'

    if os.path.exists(zip_path) and not os.path.exists(extract_path):
        print(f'Extracting {dataset_type}...')
        try:
            with zipfile.ZipFile(zip_path, 'r') as zip_ref:
                zip_ref.extractall('DIV2K/')
            print(f'✓ Extracted {dataset_type}')
        except Exception as e:
            print(f'x Failed to extract {dataset_type}: {e}')
    else:
        print(f'✓ {dataset_type} already extracted')

# Custom Dataset Class for Knowledge Distillation
class DIV2KDataset(Dataset):
    def __init__(self, hr_dir, crop_size=128, scale_factor=4, augment=True):
        self.hr_dir = hr_dir
        self.crop_size = crop_size
        self.scale_factor = scale_factor
        self.augment = augment

    # Get all image files
    self.hr_images = []
    for ext in ['*.png', '*.jpg', '*.jpeg', '*.PNG', '*.JPG', '*.JPEG']:
        self.hr_images.extend(glob.glob(os.path.join(hr_dir, ext)))

    print(f"Found {len(self.hr_images)} images in {hr_dir}")

    # Transforms
    self.to_tensor = transforms.ToTensor()
    self.to_pil = transforms.ToPILImage()

    # Augmentation transforms
    if augment:
        self.augment_transforms = transforms.Compose([
            transforms.RandomHorizontalFlip(p=0.5),
            transforms.RandomVerticalFlip(p=0.5),
            transforms.RandomRotation(degrees=90, expand=False),
        ])
    else:

```

```

        self.augment_transforms = None

    def __len__(self):
        return len(self.hr_images)

    def __getitem__(self, idx):
        # Load HR image
        hr_path = self.hr_images[idx]
        try:
            hr_img = Image.open(hr_path).convert('RGB')
        except Exception as e:
            print(f"Error loading {hr_path}: {e}")
            # Return a random valid image instead
            hr_img = Image.open(self.hr_images[0]).convert('RGB')

        # Ensure minimum size
        w, h = hr_img.size
        if w < self.crop_size or h < self.crop_size:
            # Resize to minimum required size
            scale = max(self.crop_size / w, self.crop_size / h)
            new_w, new_h = int(w * scale), int(h * scale)
            hr_img = hr_img.resize((new_w, new_h), Image.BICUBIC)

        # Random crop
        hr_img = transforms.RandomCrop(self.crop_size)(hr_img)

        # Apply augmentations
        if self.augment_transforms:
            hr_img = self.augment_transforms(hr_img)

        # Convert to tensor
        hr_tensor = self.to_tensor(hr_img)

        # Create LR image using bicubic downsampling
        lr_size = self.crop_size // self.scale_factor
        lr_img = transforms.Resize(lr_size, transforms.InterpolationMode.BICUBIC)(hr_img)
        lr_tensor = self.to_tensor(lr_img)

        return lr_tensor, hr_tensor

# Setup datasets
download_div2k_dataset()

train_dataset = DIV2KDataset('DIV2K/DIV2K_train_HR', crop_size=128, augment=True)
val_dataset = DIV2KDataset('DIV2K/DIV2K_valid_HR', crop_size=128, augment=False)

train_loader = DataLoader(train_dataset, batch_size=8, shuffle=True, num_workers=2, pin_memory=True)
val_loader = DataLoader(val_dataset, batch_size=4, shuffle=False, num_workers=2, pin_memory=True)

print(f"✓ Training dataset: {len(train_dataset)} images")
print(f"✓ Validation dataset: {len(val_dataset)} images")
print(f"✓ Training batches: {len(train_loader)}")
print(f"✓ Validation batches: {len(val_loader)}")

```

```

🔄 Starting DIV2K dataset download...
Downloading train_hr dataset...
✓ Downloaded train_hr
Downloading valid_hr dataset...
✓ Downloaded valid_hr
Extracting train_hr...
✓ Extracted train_hr
Extracting valid_hr...
✓ Extracted valid_hr
Found 800 images in DIV2K/DIV2K_train_HR
Found 100 images in DIV2K/DIV2K_valid_HR
✓ Training dataset: 800 images
✓ Validation dataset: 100 images
✓ Training batches: 100
✓ Validation batches: 25

```

**Phase 4: Enhanced 25-Epoch Training**

```

# Enhanced Knowledge Distillation Loss for 25 epochs
class EnhancedKnowledgeDistillationLoss(nn.Module):
    def __init__(self, alpha=0.7, temperature=4, perceptual_weight=0.1):
        super().__init__()
        self.alpha = alpha
        self.temperature = temperature
        self.perceptual_weight = perceptual_weight
        self.mse_loss = nn.MSELoss()
        self.l1_loss = nn.L1Loss()

    def forward(self, student_output, teacher_output, target):
        # Reconstruction loss (student vs ground truth)
        recon_loss = self.mse_loss(student_output, target)

        # Feature distillation loss (student vs teacher)
        if teacher_output.shape != student_output.shape:
            teacher_output = F.interpolate(teacher_output,
                                           size=student_output.shape[2:],
                                           mode='bilinear',
                                           align_corners=False)

        feature_loss = self.mse_loss(student_output, teacher_output)

        # Enhanced perceptual loss (edge-based)
        edge_loss = self._edge_loss(student_output, target)

        # Combined loss
        total_loss = (self.alpha * recon_loss +
                     (1 - self.alpha) * feature_loss +
                     self.perceptual_weight * edge_loss)

        return total_loss, recon_loss, feature_loss, edge_loss

    def _edge_loss(self, pred, target):
        """Calculate edge-based perceptual loss"""
        # Simple edge detection using gradients
        def get_edges(x):
            grad_x = torch.abs(x[:, :, :, :-1] - x[:, :, :, 1:])
            grad_y = torch.abs(x[:, :, :-1, :] - x[:, :, 1:, :])
            return grad_x, grad_y

        pred_grad_x, pred_grad_y = get_edges(pred)
        target_grad_x, target_grad_y = get_edges(target)

        edge_loss = (self.mse_loss(pred_grad_x, target_grad_x) +
                    self.mse_loss(pred_grad_y, target_grad_y))

        return edge_loss

# Enhanced training function for 25 epochs
def train_student_model_25_epochs(student_model, teacher_model, train_loader, val_loader, epochs=25):
    """Enhanced training with 25 epochs and improved monitoring"""

    print("Starting Enhanced Knowledge Distillation Training (25 Epochs)...")
    print(f"Student model parameters: {sum(p.numel() for p in student_model.parameters()),}")

    # Enhanced loss and optimizer
    criterion = EnhancedKnowledgeDistillationLoss(alpha=0.7, temperature=4, perceptual_weight=0.1)
    optimizer = optim.Adam(student_model.parameters(), lr=1e-4, weight_decay=1e-6)

    # Enhanced scheduler for 25 epochs
    scheduler = optim.lr_scheduler.MultiStepLR(optimizer, milestones=[10, 20], gamma=0.5)

    # Training tracking
    best_val_loss = float('inf')

```

```

best_ssim = 0.0
train_losses = []
val_losses = []
recon_losses = []
feature_losses = []
edge_losses = []
ssim_history = []

# Early stopping parameters
patience = 8
patience_counter = 0

# Training loop
for epoch in range(epochs):
    print(f"\nEpoch {epoch+1}/{epochs}")
    print("-" * 50)

    # Training phase
    student_model.train()
    epoch_train_loss = 0.0
    epoch_recon_loss = 0.0
    epoch_feature_loss = 0.0
    epoch_edge_loss = 0.0

    train_pbar = tqdm(train_loader, desc=f'Training Epoch {epoch+1}')

    for batch_idx, (lr_batch, hr_batch) in enumerate(train_pbar):
        lr_batch, hr_batch = lr_batch.to(device), hr_batch.to(device)

        # Get teacher predictions (no gradients)
        with torch.no_grad():
            teacher_outputs = []
            for i in range(lr_batch.size(0)):
                # Convert tensor to numpy for teacher model
                lr_img = lr_batch[i].cpu().numpy().transpose(1, 2, 0)
                lr_img = (lr_img * 255).astype(np.uint8)

                # Get teacher enhancement
                teacher_enhanced = teacher_model.enhance(lr_img)

                # Convert back to tensor
                teacher_enhanced = torch.from_numpy(teacher_enhanced).float() / 255.0
                teacher_enhanced = teacher_enhanced.permute(2, 0, 1)
                teacher_outputs.append(teacher_enhanced)

            teacher_batch = torch.stack(teacher_outputs).to(device)

        # Student forward pass
        optimizer.zero_grad()
        student_output = student_model(lr_batch)

        # Calculate enhanced loss
        total_loss, recon_loss, feature_loss, edge_loss = criterion(
            student_output, teacher_batch, hr_batch
        )

        # Backward pass
        total_loss.backward()
        torch.nn.utils.clip_grad_norm_(student_model.parameters(), max_norm=1.0)
        optimizer.step()

        # Update metrics
        epoch_train_loss += total_loss.item()
        epoch_recon_loss += recon_loss.item()
        epoch_feature_loss += feature_loss.item()
        epoch_edge_loss += edge_loss.item()

        # Update progress bar
        train_pbar.set_postfix({

```



```

        'Loss': f'{total_loss.item():.4f}',
        'Recon': f'{recon_loss.item():.4f}',
        'Feature': f'{feature_loss.item():.4f}',
        'Edge': f'{edge_loss.item():.4f}'
    })

# Calculate average training losses
avg_train_loss = epoch_train_loss / len(train_loader)
avg_recon_loss = epoch_recon_loss / len(train_loader)
avg_feature_loss = epoch_feature_loss / len(train_loader)
avg_edge_loss = epoch_edge_loss / len(train_loader)

# Validation phase
student_model.eval()
epoch_val_loss = 0.0
val_ssim_scores = []

val_pbar = tqdm(val_loader, desc=f'Validation Epoch {epoch+1}')

with torch.no_grad():
    for lr_batch, hr_batch in val_pbar:
        lr_batch, hr_batch = lr_batch.to(device), hr_batch.to(device)

        # Student prediction
        student_output = student_model(lr_batch)

        # Validation loss (reconstruction only)
        val_loss = F.mse_loss(student_output, hr_batch)
        epoch_val_loss += val_loss.item()

        # Calculate SSIM for first image in batch
        student_img = student_output[0].cpu().numpy().transpose(1, 2, 0)
        hr_img = hr_batch[0].cpu().numpy().transpose(1, 2, 0)

        student_img = np.clip(student_img * 255, 0, 255).astype(np.uint8)
        hr_img = np.clip(hr_img * 255, 0, 255).astype(np.uint8)

        # Convert to grayscale for SSIM
        student_gray = cv2.cvtColor(student_img, cv2.COLOR_RGB2GRAY)
        hr_gray = cv2.cvtColor(hr_img, cv2.COLOR_RGB2GRAY)

        ssim_score = ssim(hr_gray, student_gray, data_range=255)
        val_ssim_scores.append(ssim_score)

    val_pbar.set_postfix({
        'Val Loss': f'{val_loss.item():.4f}',
        'SSIM': f'{ssim_score:.4f}'
    })

avg_val_loss = epoch_val_loss / len(val_loader)
avg_ssim = np.mean(val_ssim_scores)

# Store metrics
train_losses.append(avg_train_loss)
val_losses.append(avg_val_loss)
recon_losses.append(avg_recon_loss)
feature_losses.append(avg_feature_loss)
edge_losses.append(avg_edge_loss)
ssim_history.append(avg_ssim)

# Print epoch summary
print(f"\nEpoch {epoch+1} Summary:")
print(f"  Train Loss: {avg_train_loss:.4f}")
print(f"  Val Loss: {avg_val_loss:.4f}")
print(f"  Reconstruction Loss: {avg_recon_loss:.4f}")
print(f"  Feature Loss: {avg_feature_loss:.4f}")
print(f"  Edge Loss: {avg_edge_loss:.4f}")
print(f"  Validation SSIM: {avg_ssim:.4f} ({avg_ssim*100:.2f}%)")
print(f"  Learning Rate: {optimizer.param_groups[0]['lr']:.6f}")

```

```

# Enhanced model saving
if avg_val_loss < best_val_loss or avg_ssim > best_ssim:
    if avg_val_loss < best_val_loss:
        best_val_loss = avg_val_loss
        patience_counter = 0
    if avg_ssim > best_ssim:
        best_ssim = avg_ssim

    torch.save({
        'epoch': epoch,
        'model_state_dict': student_model.state_dict(),
        'optimizer_state_dict': optimizer.state_dict(),
        'val_loss': avg_val_loss,
        'ssim': avg_ssim,
        'train_losses': train_losses,
        'val_losses': val_losses,
        'ssim_history': ssim_history
    }, 'best_student_model_25epochs.pth')
    print(f" ✓ New best model saved! (Val Loss: {avg_val_loss:.4f}, SSIM: {avg_ssim:.4f})")
else:
    patience_counter += 1

# Early stopping check
if patience_counter >= patience:
    print(f"\nEarly stopping triggered after {epoch+1} epochs (patience: {patience})")
    break

# Update learning rate
scheduler.step()

# Plot progress every 5 epochs
if (epoch + 1) % 5 == 0:
    plot_enhanced_training_progress(train_losses, val_losses, recon_losses,
                                    feature_losses, edge_losses, ssim_history)

print(f"\nTraining completed!")
print(f"Best validation loss: {best_val_loss:.4f}")
print(f"Best SSIM score: {best_ssim:.4f} ({best_ssim*100:.2f}%)")

return train_losses, val_losses, recon_losses, feature_losses, edge_losses, ssim_history

# Enhanced training progress visualization
def plot_enhanced_training_progress(train_losses, val_losses, recon_losses,
                                    feature_losses, edge_losses, ssim_history):
    """Plot enhanced training progress with all metrics"""
    fig, axes = plt.subplots(2, 3, figsize=(18, 10))

    # Overall losses
    axes[0, 0].plot(train_losses, label='Training Loss', color='blue', marker='o')
    axes[0, 0].plot(val_losses, label='Validation Loss', color='red', marker='s')
    axes[0, 0].set_xlabel('Epoch')
    axes[0, 0].set_ylabel('Loss')
    axes[0, 0].set_title('Overall Training Progress (25 Epochs)')
    axes[0, 0].legend()
    axes[0, 0].grid(True)

    # Loss components
    axes[0, 1].plot(recon_losses, label='Reconstruction Loss', color='green', marker='^')
    axes[0, 1].plot(feature_losses, label='Feature Distillation Loss', color='orange', marker='v')
    axes[0, 1].plot(edge_losses, label='Edge Loss', color='purple', marker='d')
    axes[0, 1].set_xlabel('Epoch')
    axes[0, 1].set_ylabel('Loss')
    axes[0, 1].set_title('Loss Components')
    axes[0, 1].legend()
    axes[0, 1].grid(True)

    # SSIM Progress
    axes[0, 2].plot(ssim_history, label='Validation SSIM', color='cyan', marker='*')

```

```

axes[0, 2].axhline(y=0.9, color='red', linestyle='--', label='Target (90%)')
axes[0, 2].set_xlabel('Epoch')
axes[0, 2].set_ylabel('SSIM Score')
axes[0, 2].set_title('SSIM Progress')
axes[0, 2].legend()
axes[0, 2].grid(True)

# Loss ratio analysis
if len(recon_losses) > 0 and len(feature_losses) > 0:
    ratios = [r/f if f > 0 else 0 for r, f in zip(recon_losses, feature_losses)]
    axes[1, 0].plot(ratios, label='Recon/Feature Ratio', color='purple', marker='d')
    axes[1, 0].set_xlabel('Epoch')
    axes[1, 0].set_ylabel('Ratio')
    axes[1, 0].set_title('Loss Ratio Analysis')
    axes[1, 0].legend()
    axes[1, 0].grid(True)

# Validation trend
if len(val_losses) > 1:
    val_trend = np.diff(val_losses)
    axes[1, 1].plot(val_trend, label='Validation Trend', color='red', marker='x')
    axes[1, 1].axhline(y=0, color='black', linestyle='--', alpha=0.5)
    axes[1, 1].set_xlabel('Epoch')
    axes[1, 1].set_ylabel('Loss Change')
    axes[1, 1].set_title('Validation Loss Trend')
    axes[1, 1].legend()
    axes[1, 1].grid(True)

# SSIM improvement
if len(ssim_history) > 1:
    ssim_trend = np.diff(ssim_history)
    axes[1, 2].plot(ssim_trend, label='SSIM Improvement', color='green', marker='+')
    axes[1, 2].axhline(y=0, color='black', linestyle='--', alpha=0.5)
    axes[1, 2].set_xlabel('Epoch')
    axes[1, 2].set_ylabel('SSIM Change')
    axes[1, 2].set_title('SSIM Improvement Trend')
    axes[1, 2].legend()
    axes[1, 2].grid(True)

plt.tight_layout()
plt.show()

# Load best model function
def load_best_student_model_25epochs(student_model, checkpoint_path='best_student_model_25epochs.pth'):
    """Load the best trained student model from 25 epochs"""
    if os.path.exists(checkpoint_path):
        try:
            checkpoint = torch.load(checkpoint_path, map_location=device, weights_only=False)
            student_model.load_state_dict(checkpoint['model_state_dict'])
            print(f"✓ Loaded best model from epoch {checkpoint['epoch']}")
            print(f" Validation Loss: {checkpoint['val_loss']:.4f}")
            print(f" SSIM Score: {checkpoint['ssim']:.4f} ({checkpoint['ssim']*100:.2f}%)")
            return checkpoint
        except Exception as e:
            print(f"x Error loading checkpoint: {e}")
            return None
    else:
        print(f"x Checkpoint not found: {checkpoint_path}")
        return None

print("Enhanced 25-epoch training setup complete!")
print("Run the following to start training:")
print("results = train_student_model_25_epochs(student_model, teacher_model, train_loader, val_loader, epochs=25)")

```

➡ Enhanced 25-epoch training setup complete!  
Run the following to start training:  
results = train\_student\_model\_25\_epochs(student\_model, teacher\_model, train\_loader, val\_loader, epochs=25)

**Running the above command**

```
results = train_student_model_25_epochs(student_model, teacher_model, train_loader, val_loader, epochs=25)
```

Starting Enhanced Knowledge Distillation Training (25 Epochs)...  
Student model parameters: 944,323

Epoch 1/25

-----  
Training Epoch 1: 100%|██████████| 100/100 [06:00<00:00, 3.60s/it, Loss=0.0295, Recon=0.0196, Feature=0.0493,  
Validation Epoch 1: 100%|██████████| 25/25 [00:15<00:00, 1.65it/s, Val Loss=0.0215, SSIM=0.5924]

Epoch 1 Summary:

Train Loss: 0.0349  
Val Loss: 0.0170  
Reconstruction Loss: 0.0299  
Feature Loss: 0.0428  
Edge Loss: 0.0110  
Validation SSIM: 0.5100 (51.00%)  
Learning Rate: 0.000100  
✓ New best model saved! (Val Loss: 0.0170, SSIM: 0.5100)

Epoch 2/25

-----  
Training Epoch 2: 100%|██████████| 100/100 [05:54<00:00, 3.55s/it, Loss=0.0256, Recon=0.0161, Feature=0.0449,  
Validation Epoch 2: 100%|██████████| 25/25 [00:14<00:00, 1.77it/s, Val Loss=0.0123, SSIM=0.8320]

Epoch 2 Summary:

Train Loss: 0.0283  
Val Loss: 0.0157  
Reconstruction Loss: 0.0190  
Feature Loss: 0.0465  
Edge Loss: 0.0108  
Validation SSIM: 0.5343 (53.43%)  
Learning Rate: 0.000100  
✓ New best model saved! (Val Loss: 0.0157, SSIM: 0.5343)

Epoch 3/25

-----  
Training Epoch 3: 100%|██████████| 100/100 [06:03<00:00, 3.63s/it, Loss=0.0294, Recon=0.0192, Feature=0.0497,  
Validation Epoch 3: 100%|██████████| 25/25 [00:14<00:00, 1.77it/s, Val Loss=0.0106, SSIM=0.8768]

Epoch 3 Summary:

Train Loss: 0.0271  
Val Loss: 0.0136  
Reconstruction Loss: 0.0171  
Feature Loss: 0.0468  
Edge Loss: 0.0104  
Validation SSIM: 0.5942 (59.42%)  
Learning Rate: 0.000100  
✓ New best model saved! (Val Loss: 0.0136, SSIM: 0.5942)

Epoch 4/25

-----  
Training Epoch 4: 100%|██████████| 100/100 [06:05<00:00, 3.65s/it, Loss=0.0235, Recon=0.0151, Feature=0.0400,  
Validation Epoch 4: 100%|██████████| 25/25 [00:14<00:00, 1.71it/s, Val Loss=0.0128, SSIM=0.9087]

Epoch 4 Summary:

Train Loss: 0.0268  
Val Loss: 0.0120  
Reconstruction Loss: 0.0161  
Feature Loss: 0.0483  
Edge Loss: 0.0107  
Validation SSIM: 0.5981 (59.81%)  
Learning Rate: 0.000100  
✓ New best model saved! (Val Loss: 0.0120, SSIM: 0.5981)

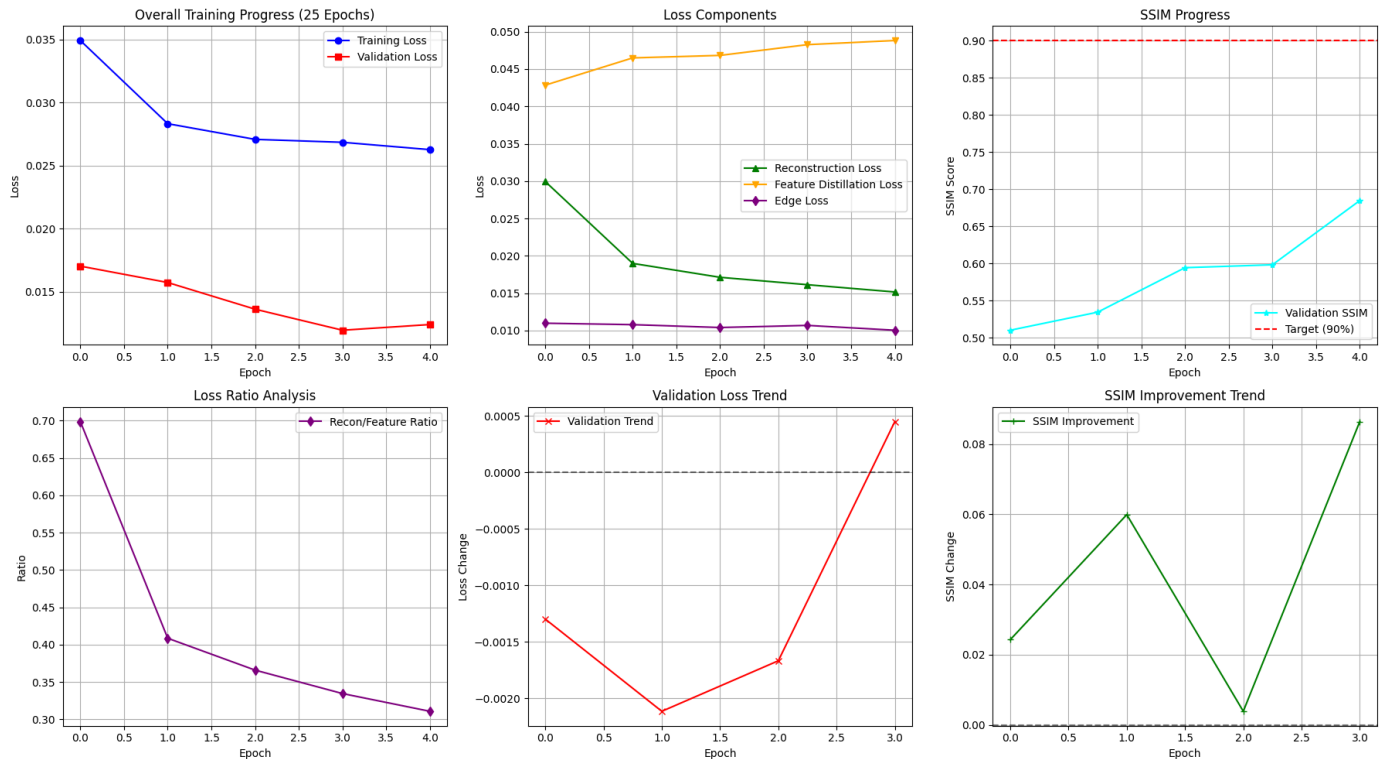
Epoch 5/25

-----  
Training Epoch 5: 100%|██████████| 100/100 [06:02<00:00, 3.63s/it, Loss=0.0200, Recon=0.0117, Feature=0.0371,  
Validation Epoch 5: 100%|██████████| 25/25 [00:14<00:00, 1.76it/s, Val Loss=0.0145, SSIM=0.9203]

Epoch 5 Summary:

Train Loss: 0.0263  
Val Loss: 0.0124  
Reconstruction Loss: 0.0152  
Feature Loss: 0.0488  
Edge Loss: 0.0101  
Validation SSIM: 0.6845 (68.45%)  
Learning Rate: 0.000100

✓ New best model saved! (Val Loss: 0.0124, SSIM: 0.6845)



Epoch 6/25

Training Epoch 6: 100% | 100/100 [05:58<00:00, 3.59s/it, Loss=0.0279, Recon=0.0142, Feature=0.0568, Validation Epoch 6: 100% | 25/25 [00:13<00:00, 1.79it/s, Val Loss=0.0075, SSIM=0.6440]

Epoch 6 Summary:

Train Loss: 0.0259  
Val Loss: 0.0117  
Reconstruction Loss: 0.0149  
Feature Loss: 0.0481  
Edge Loss: 0.0105  
Validation SSIM: 0.6760 (67.60%)  
Learning Rate: 0.000100  
✓ New best model saved! (Val Loss: 0.0117, SSIM: 0.6760)

Epoch 7/25

Training Epoch 7: 100% | 100/100 [05:58<00:00, 3.59s/it, Loss=0.0202, Recon=0.0105, Feature=0.0407, Validation Epoch 7: 100% | 25/25 [00:13<00:00, 1.79it/s, Val Loss=0.0113, SSIM=0.9421]

Epoch 7 Summary:

Train Loss: 0.0260  
Val Loss: 0.0120  
Reconstruction Loss: 0.0145  
Feature Loss: 0.0492  
Edge Loss: 0.0100  
Validation SSIM: 0.6924 (69.24%)  
Learning Rate: 0.000100  
✓ New best model saved! (Val Loss: 0.0120, SSIM: 0.6924)

Epoch 8/25

Training Epoch 8: 100% | 100/100 [06:00<00:00, 3.61s/it, Loss=0.0242, Recon=0.0135, Feature=0.0462, Validation Epoch 8: 100% | 25/25 [00:14<00:00, 1.78it/s, Val Loss=0.0136, SSIM=0.7050]

Epoch 8 Summary:

Train Loss: 0.0253  
Val Loss: 0.0115  
Reconstruction Loss: 0.0139  
Feature Loss: 0.0487  
Edge Loss: 0.0095  
Validation SSIM: 0.6792 (67.92%)  
Learning Rate: 0.000100  
✓ New best model saved! (Val Loss: 0.0115, SSIM: 0.6792)

Epoch 9/25

Training Epoch 9: 100% |██████████| 100/100 [05:50<00:00, 3.50s/it, Loss=0.0277, Recon=0.0162, Feature=0.0517,  
Validation Epoch 9: 100% |██████████| 25/25 [00:21<00:00, 1.14it/s, Val Loss=0.0085, SSIM=0.9523]

Epoch 9 Summary:

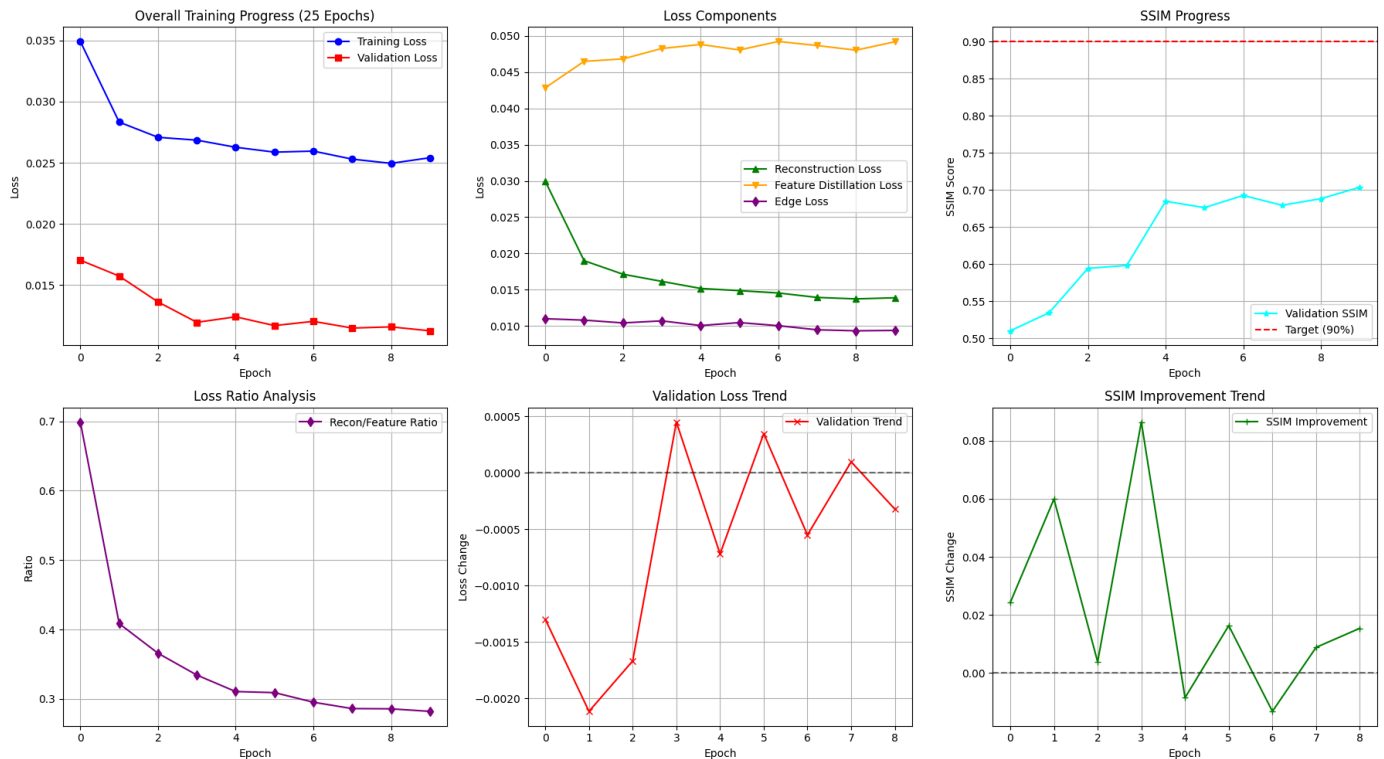
Train Loss: 0.0250  
Val Loss: 0.0116  
Reconstruction Loss: 0.0137  
Feature Loss: 0.0480  
Edge Loss: 0.0093  
Validation SSIM: 0.6881 (68.81%)  
Learning Rate: 0.000100

Epoch 10/25

Training Epoch 10: 100% |██████████| 100/100 [05:52<00:00, 3.53s/it, Loss=0.0274, Recon=0.0132, Feature=0.0590,  
Validation Epoch 10: 100% |██████████| 25/25 [00:13<00:00, 1.79it/s, Val Loss=0.0103, SSIM=0.9550]

Epoch 10 Summary:

Train Loss: 0.0254  
Val Loss: 0.0113  
Reconstruction Loss: 0.0139  
Feature Loss: 0.0492  
Edge Loss: 0.0094  
Validation SSIM: 0.7035 (70.35%)  
Learning Rate: 0.000100  
✓ New best model saved! (Val Loss: 0.0113, SSIM: 0.7035)



Epoch 11/25

Training Epoch 11: 100% |██████████| 100/100 [05:57<00:00, 3.57s/it, Loss=0.0235, Recon=0.0133, Feature=0.0443,  
Validation Epoch 11: 100% |██████████| 25/25 [00:14<00:00, 1.76it/s, Val Loss=0.0094, SSIM=0.9565]

Epoch 11 Summary:

Train Loss: 0.0251  
Val Loss: 0.0118  
Reconstruction Loss: 0.0137  
Feature Loss: 0.0485  
Edge Loss: 0.0097  
Validation SSIM: 0.7085 (70.85%)  
Learning Rate: 0.000050  
✓ New best model saved! (Val Loss: 0.0118, SSIM: 0.7085)

Epoch 12/25

Training Epoch 12: 100% |██████████| 100/100 [05:54<00:00, 3.55s/it, Loss=0.0226, Recon=0.0113, Feature=0.0471,

Epoch 12 Summary:  
Train Loss: 0.0254  
Val Loss: 0.0116  
Reconstruction Loss: 0.0138  
Feature Loss: 0.0493  
Edge Loss: 0.0094  
Validation SSIM: 0.7492 (74.92%)  
Learning Rate: 0.000050  
✓ New best model saved! (Val Loss: 0.0116, SSIM: 0.7492)

Epoch 13/25

---

Training Epoch 13: 100%|██████████| 100/100 [05:57<00:00, 3.57s/it, Loss=0.0249, Recon=0.0137, Feature=0.0468, Validation Epoch 13: 100%|██████████| 25/25 [00:15<00:00, 1.61it/s, Val Loss=0.0105, SSIM=0.9532]

Epoch 13 Summary:  
Train Loss: 0.0250  
Val Loss: 0.0110  
Reconstruction Loss: 0.0135  
Feature Loss: 0.0486  
Edge Loss: 0.0095  
Validation SSIM: 0.7043 (70.43%)  
Learning Rate: 0.000050  
✓ New best model saved! (Val Loss: 0.0110, SSIM: 0.7043)

Epoch 14/25

---

Training Epoch 14: 100%|██████████| 100/100 [05:59<00:00, 3.59s/it, Loss=0.0289, Recon=0.0147, Feature=0.0591, Validation Epoch 14: 100%|██████████| 25/25 [00:15<00:00, 1.62it/s, Val Loss=0.0098, SSIM=0.9588]

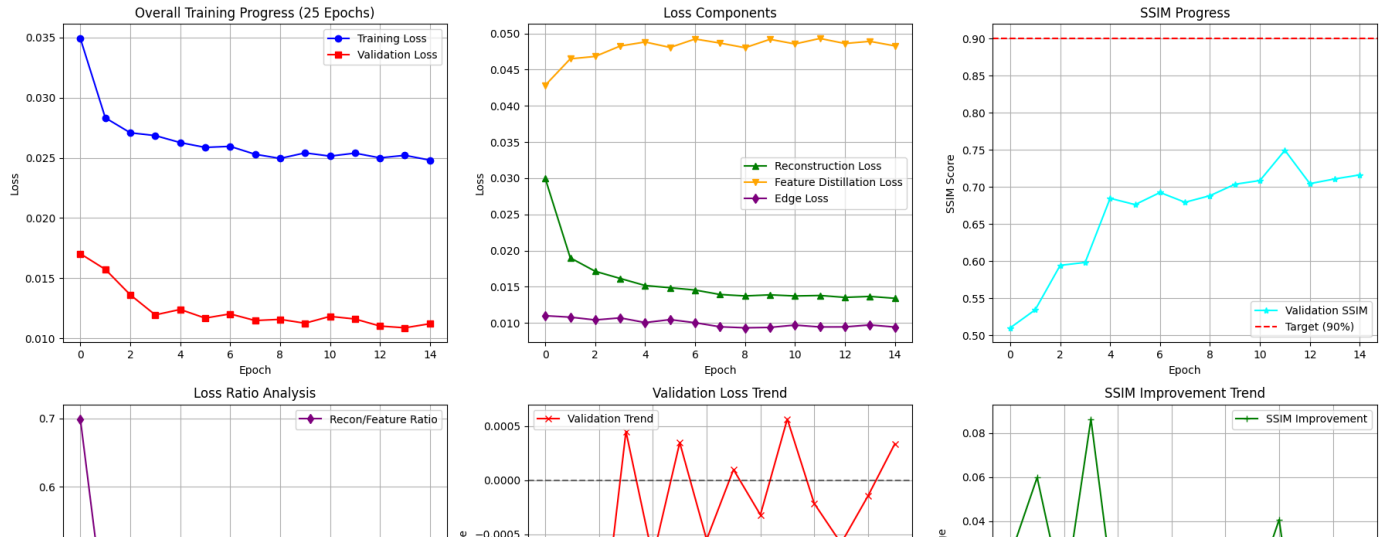
Epoch 14 Summary:  
Train Loss: 0.0252  
Val Loss: 0.0109  
Reconstruction Loss: 0.0137  
Feature Loss: 0.0489  
Edge Loss: 0.0097  
Validation SSIM: 0.7107 (71.07%)  
Learning Rate: 0.000050  
✓ New best model saved! (Val Loss: 0.0109, SSIM: 0.7107)

Epoch 15/25




---

Training Epoch 15: 100%|██████████| 100/100 [05:51<00:00, 3.52s/it, Loss=0.0242, Recon=0.0143, Feature=0.0432, Validation Epoch 15: 100%|██████████| 25/25 [00:14<00:00, 1.76it/s, Val Loss=0.0095, SSIM=0.9584]

Epoch 15 Summary:  
Train Loss: 0.0248  
Val Loss: 0.0112  
Reconstruction Loss: 0.0134  
Feature Loss: 0.0483  
Edge Loss: 0.0094  
Validation SSIM: 0.7161 (71.61%)  
Learning Rate: 0.000050







## Epoch 16/25

Training Epoch 16: 100%|██████████| 100/100 [06:01<00:00, 3.61s/it, Loss=0.0254, Recon=0.0121, Feature=0.0544, Validation Epoch 16: 100%|██████████| 25/25 [00:14<00:00, 1.74it/s, Val Loss=0.0098, SSIM=0.6952]

## Epoch 16 Summary:

Train Loss: 0.0254  
 Val Loss: 0.0101  
 Reconstruction Loss: 0.0138  
 Feature Loss: 0.0493  
 Edge Loss: 0.0096  
 Validation SSIM: 0.7167 (71.67%)  
 Learning Rate: 0.000050  
 ✓ New best model saved! (Val Loss: 0.0101, SSIM: 0.7167)

## Epoch 17/25

Training Epoch 17: 100%|██████████| 100/100 [06:00<00:00, 3.61s/it, Loss=0.0244, Recon=0.0125, Feature=0.0504, Validation Epoch 17: 100%|██████████| 25/25 [00:14<00:00, 1.70it/s, Val Loss=0.0112, SSIM=0.9633]

## Epoch 17 Summary:

Train Loss: 0.0255  
 Val Loss: 0.0112  
 Reconstruction Loss: 0.0138  
 Feature Loss: 0.0492  
 Edge Loss: 0.0101  
 Validation SSIM: 0.7279 (72.79%)  
 Learning Rate: 0.000050

## Epoch 18/25

Training Epoch 18: 100%|██████████| 100/100 [06:10<00:00, 3.70s/it, Loss=0.0251, Recon=0.0122, Feature=0.0534, Validation Epoch 18: 100%|██████████| 25/25 [00:15<00:00, 1.61it/s, Val Loss=0.0092, SSIM=0.5831]

## Epoch 18 Summary:

Train Loss: 0.0250  
 Val Loss: 0.0103  
 Reconstruction Loss: 0.0133  
 Feature Loss: 0.0493  
 Edge Loss: 0.0092  
 Validation SSIM: 0.7312 (73.12%)  
 Learning Rate: 0.000050

## Epoch 19/25

Training Epoch 19: 100%|██████████| 100/100 [06:01<00:00, 3.61s/it, Loss=0.0242, Recon=0.0128, Feature=0.0467, Validation Epoch 19: 100%|██████████| 25/25 [00:14<00:00, 1.71it/s, Val Loss=0.0091, SSIM=0.9582]

## Epoch 19 Summary:

Train Loss: 0.0252  
 Val Loss: 0.0114  
 Reconstruction Loss: 0.0135  
 Feature Loss: 0.0496  
 Edge Loss: 0.0092  
 Validation SSIM: 0.7104 (71.04%)  
 Learning Rate: 0.000050

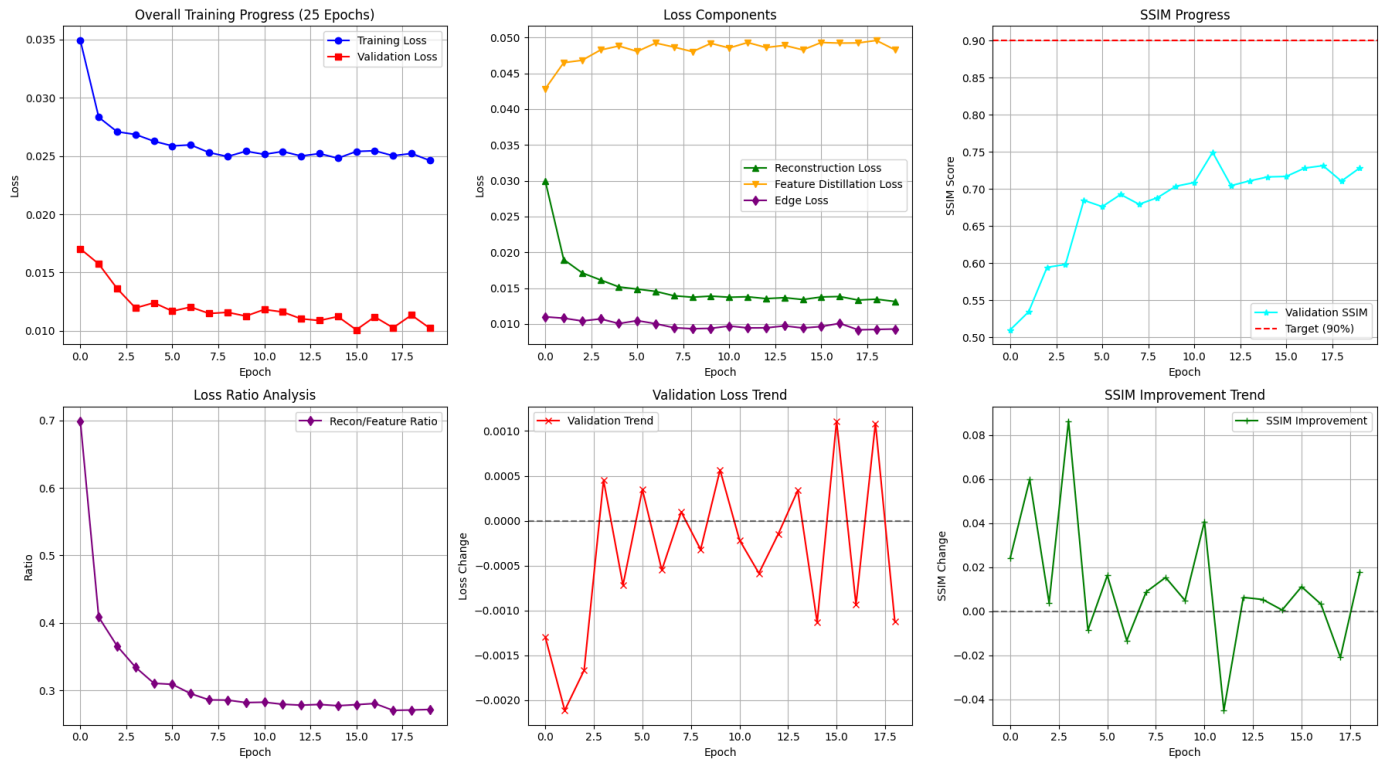
## Epoch 20/25

Training Epoch 20: 100%|██████████| 100/100 [05:57<00:00, 3.57s/it, Loss=0.0257, Recon=0.0137, Feature=0.0500, Validation Epoch 20: 100%|██████████| 25/25 [00:14<00:00, 1.68it/s, Val Loss=0.0093, SSIM=0.9567]

## Epoch 20 Summary:

Train Loss: 0.0246  
 Val Loss: 0.0102  
 Reconstruction Loss: 0.0131  
 Feature Loss: 0.0483

Edge Loss: 0.0093  
 Validation SSIM: 0.7282 (72.82%)  
 Learning Rate: 0.000050



Epoch 21/25

Training Epoch 21: 100% |██████████| 100/100 [06:03<00:00, 3.64s/it, Loss=0.0282, Recon=0.0139, Feature=0.0596,  
 Validation Epoch 21: 100% |██████████| 25/25 [00:14<00:00, 1.73it/s, Val Loss=0.0087, SSIM=0.9619]

Epoch 21 Summary:

Train Loss: 0.0250  
 Val Loss: 0.0112  
 Reconstruction Loss: 0.0132  
 Feature Loss: 0.0493  
 Edge Loss: 0.0091  
 Validation SSIM: 0.7067 (70.67%)  
 Learning Rate: 0.000025

Epoch 22/25

Training Epoch 22: 100% |██████████| 100/100 [05:56<00:00, 3.57s/it, Loss=0.0189, Recon=0.0111, Feature=0.0333,  
 Validation Epoch 22: 100% |██████████| 25/25 [00:14<00:00, 1.68it/s, Val Loss=0.0094, SSIM=0.9619]

Epoch 22 Summary:

Train Loss: 0.0250  
 Val Loss: 0.0114  
 Reconstruction Loss: 0.0134  
 Feature Loss: 0.0492  
 Edge Loss: 0.0094  
 Validation SSIM: 0.7577 (75.77%)  
 Learning Rate: 0.000025  
 ✓ New best model saved! (Val Loss: 0.0114, SSIM: 0.7577)

Epoch 23/25

Training Epoch 23: 100% |██████████| 100/100 [06:05<00:00, 3.66s/it, Loss=0.0277, Recon=0.0142, Feature=0.0563,  
 Validation Epoch 23: 100% |██████████| 25/25 [00:14<00:00, 1.70it/s, Val Loss=0.0109, SSIM=0.9583]

Epoch 23 Summary:

Train Loss: 0.0253  
 Val Loss: 0.0111  
 Reconstruction Loss: 0.0135  
 Feature Loss: 0.0493  
 Edge Loss: 0.0099  
 Validation SSIM: 0.7077 (70.77%)  
 Learning Rate: 0.000025

Epoch 24/25

-----  
Training Epoch 24: 100%|██████████| 100/100 [06:00<00:00, 3.61s/it, Loss=0.0260, Recon=0.0139, Feature=0.0511,  
Validation Epoch 24: 100%|██████████| 25/25 [00:14<00:00, 1.72it/s, Val Loss=0.0094, SSIM=0.9610]

Epoch 24 Summary:

Train Loss: 0.0255  
Val Loss: 0.0106  
Reconstruction Loss: 0.0135  
Feature Loss: 0.0502  
Edge Loss: 0.0094  
Validation SSIM: 0.6917 (69.17%)  
Learning Rate: 0.000025

Epoch 25/25

-----  
Training Epoch 25: 100%|██████████| 100/100 [05:55<00:00, 3.56s/it, Loss=0.0210, Recon=0.0112, Feature=0.0413,  
Validation Epoch 25: 100%|██████████| 25/25 [00:14<00:00, 1.72it/s, Val Loss=0.0099, SSIM=0.9631]

Epoch 25 Summary:

Train Loss: 0.0247  
Val Loss: 0.0109  
Reconstruction Loss: 0.0131  
Feature Loss: 0.0488  
Edge Loss: 0.0090  
Validation SSIM: 0.7157 (71.57%)  
Learning Rate: 0.000025

Early stopping triggered after 25 epochs (patience: 8)

Training completed!

Best validation loss: 0.0101

Best SSIM score: 0.7577 (75.77%)



## Downloading the best Model into Local Machine

```
# Download the best model file to your local machine
from google.colab import files

# Check if the model file exists and download it
import os

model_files = [
    'best_student_model.pth',
    'best_student_model_25epochs.pth' # If you used 25 epochs
]

for model_file in model_files:
    if os.path.exists(model_file):
        print(f"Downloading {model_file}...")
        files.download(model_file)
        print(f"✓ {model_file} downloaded successfully!")
        break
    else:
        print("❌ No model checkpoint found!")
        # List all .pth files in current directory
        pth_files = [f for f in os.listdir('.') if f.endswith('.pth')]
        if pth_files:
            print("Available .pth files:")
            for f in pth_files:
                print(f" • {f}")
```

📁 Downloading best\_student\_model\_25epochs.pth...  
 ✓ best\_student\_model\_25epochs.pth downloaded successfully!

## Phase 5: Complete Video Processing Pipeline Setup

```
# Import required libraries
import os
import cv2
import numpy as np
import base64
from IPython.display import display, HTML
from google.colab import files
from tqdm import tqdm
from skimage.metrics import structural_similarity as ssim
import matplotlib.pyplot as plt
import time

def upload_video_file():
    """Upload video file from local machine"""
    print("Please upload a video file...")
    print("Supported formats: .mp4, .avi, .mov, .mkv")

    uploaded = files.upload()

    if not uploaded:
        print("No file uploaded!")
        return None

    video_path = list(uploaded.keys())[0]
    print(f"✓ Video uploaded: {video_path}")
    return video_path

def extract_video_info(video_path):
    """Extract basic video information"""
    cap = cv2.VideoCapture(video_path)
    if not cap.isOpened():
        return None
```

```

    fps = int(cap.get(cv2.CAP_PROP_FPS))
    width = int(cap.get(cv2.CAP_PROP_FRAME_WIDTH))
    height = int(cap.get(cv2.CAP_PROP_FRAME_HEIGHT))
    frame_count = int(cap.get(cv2.CAP_PROP_FRAME_COUNT))
    duration = frame_count / fps if fps > 0 else 0

    cap.release()

    return {
        'fps': fps,
        'width': width,
        'height': height,
        'frame_count': frame_count,
        'duration': duration
    }

def simulate_network_blur(frame, blur_factor=4):
    """Simulate poor network conditions by downscaling and upscaling"""
    h, w = frame.shape[:2]
    small_frame = cv2.resize(frame, (w//blur_factor, h//blur_factor), interpolation=cv2.INTER_CUBIC)
    blurred_frame = cv2.resize(small_frame, (w, h), interpolation=cv2.INTER_CUBIC)
    return blurred_frame

def enhance_frame_with_student(frame, student_model):
    """Enhance a single frame using the trained student model"""
    try:
        frame_rgb = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)
        frame_tensor = torch.from_numpy(frame_rgb).float() / 255.0
        frame_tensor = frame_tensor.permute(2, 0, 1).unsqueeze(0).to(device)

        student_model.eval()
        with torch.no_grad():
            enhanced_tensor = student_model(frame_tensor)

        enhanced_frame = enhanced_tensor.squeeze(0).permute(1, 2, 0).cpu().numpy()
        enhanced_frame = np.clip(enhanced_frame * 255, 0, 255).astype(np.uint8)
        enhanced_frame = cv2.cvtColor(enhanced_frame, cv2.COLOR_RGB2BGR)

        if enhanced_frame.shape[:2] != frame.shape[:2]:
            enhanced_frame = cv2.resize(enhanced_frame, (frame.shape[1], frame.shape[0]))

        return enhanced_frame

    except Exception as e:
        print(f"Enhancement failed: {e}")
        return frame

def process_video_frames(video_path, student_model, max_frames=None):
    """Process video frames: extract, blur, enhance"""
    print(f"Processing video: {video_path}")

    video_info = extract_video_info(video_path)
    if video_info is None:
        print("x Failed to read video file")
        return None, None, None, None

    cap = cv2.VideoCapture(video_path)
    original_frames = []
    blurred_frames = []
    enhanced_frames = []
    frame_metrics = []

    frame_count = 0
    max_process = max_frames if max_frames else video_info['frame_count']

    pbar = tqdm(total=min(max_process, video_info['frame_count']), desc="Processing frames")

    while True:
        ret, frame = cap.read()

```

```

    if not ret or frame_count >= max_process:
        break

    original_frames.append(frame.copy())
    blurred_frame = simulate_network_blur(frame, blur_factor=4)
    blurred_frames.append(blurred_frame)

    start_time = time.time()
    enhanced_frame = enhance_frame_with_student(blurred_frame, student_model)
    processing_time = time.time() - start_time

    enhanced_frames.append(enhanced_frame)

    frame_ssim = ssim(cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY),
                      cv2.cvtColor(enhanced_frame, cv2.COLOR_BGR2GRAY), data_range=255)

    frame_metrics.append({
        'frame_id': frame_count,
        'ssim': frame_ssim,
        'processing_time': processing_time,
        'fps': 1.0 / processing_time if processing_time > 0 else 0
    })

    frame_count += 1
    pbar.update(1)
    pbar.set_postfix({'SSIM': f'{frame_ssim:.3f}'})

cap.release()
pbar.close()

print(f"✓ Processed {len(original_frames)} frames")
return original_frames, blurred_frames, enhanced_frames, frame_metrics

def create_output_videos(original_frames, blurred_frames, enhanced_frames, video_info, output_dir='output_videos'):
    """Create output videos from processed frames"""
    if not os.path.exists(output_dir):
        os.makedirs(output_dir)

    fps = video_info['fps']
    width, height = video_info['width'], video_info['height']
    fourcc = cv2.VideoWriter_fourcc(*'mp4v')

    videos_created = []

    # Enhanced video
    enhanced_path = os.path.join(output_dir, 'enhanced_video.mp4')
    out_enhanced = cv2.VideoWriter(enhanced_path, fourcc, fps, (width, height))

    print("Creating enhanced video...")
    for frame in tqdm(enhanced_frames, desc="Enhanced video"):
        if frame.shape[:2] != (height, width):
            frame = cv2.resize(frame, (width, height))
        out_enhanced.write(frame)

    out_enhanced.release()
    videos_created.append(enhanced_path)
    print(f"✓ Enhanced video saved: {enhanced_path}")

    # Blurred video
    blurred_path = os.path.join(output_dir, 'blurred_video.mp4')
    out_blurred = cv2.VideoWriter(blurred_path, fourcc, fps, (width, height))

    print("Creating blurred video...")
    for frame in tqdm(blurred_frames, desc="Blurred video"):
        if frame.shape[:2] != (height, width):
            frame = cv2.resize(frame, (width, height))
        out_blurred.write(frame)

    out_blurred.release()

```

```

        videos_created.append(blurred_path)
        print(f"✓ Blurred video saved: {blurred_path}")

    return videos_created

def create_side_by_side_comparison_video(original_frames, blurred_frames, enhanced_frames, video_info, output_path=
    """Create side-by-side comparison video with labels"""
    fps = video_info['fps']
    width, height = video_info['width'], video_info['height']
    comparison_width = width * 3
    fourcc = cv2.VideoWriter_fourcc(*'mp4v')
    out = cv2.VideoWriter(output_path, fourcc, fps, (comparison_width, height))

    print(f"Creating comparison video: {output_path}")

    for i, (orig, blur, enh) in enumerate(tqdm(zip(original_frames, blurred_frames, enhanced_frames),
        desc="Creating comparison", total=len(original_frames))):

        orig_resized = cv2.resize(orig, (width, height))
        blur_resized = cv2.resize(blur, (width, height))
        enh_resized = cv2.resize(enh, (width, height))

        font = cv2.FONT_HERSHEY_SIMPLEX
        cv2.putText(orig_resized, 'GROUND TRUTH', (10, 40), font, 1.2, (0, 255, 0), 2)
        cv2.putText(blur_resized, 'BLURRED', (10, 40), font, 1.2, (0, 0, 255), 2)
        cv2.putText(enh_resized, 'ENHANCED', (10, 40), font, 1.2, (255, 0, 0), 2)

        # Add SSIM score
        frame_ssim = ssim(cv2.cvtColor(orig_resized, cv2.COLOR_BGR2GRAY),
            cv2.cvtColor(enh_resized, cv2.COLOR_BGR2GRAY), data_range=255)
        cv2.putText(enh_resized, f'SSIM: {frame_ssim:.3f}', (10, height-60), font, 0.8, (255, 255, 0), 2)

        comparison_frame = np.hstack([orig_resized, blur_resized, enh_resized])
        cv2.line(comparison_frame, (width, 0), (width, height), (255, 255, 255), 2)
        cv2.line(comparison_frame, (width*2, 0), (width*2, height), (255, 255, 255), 2)

        out.write(comparison_frame)

    out.release()
    print(f"✓ Comparison video created: {output_path}")
    return output_path

def display_video_in_colab(video_path, title="Video", width=640, height=480):
    """Display video in Google Colab using HTML5 video player"""
    try:
        with open(video_path, 'rb') as f:
            video_data = f.read()

        video_base64 = base64.b64encode(video_data).decode()

        video_html = f"""
        <div style="text-align: center; margin: 20px;">
            <h3>{title}</h3>
            <video width="{width}" height="{height}" controls>
                <source src="data:video/mp4;base64,{video_base64}" type="video/mp4">
                Your browser does not support the video tag.
            </video>
        </div>
        """

        display(HTML(video_html))
        return True

    except Exception as e:
        print(f"Error displaying video {video_path}: {e}")
        return False

def download_videos_to_local(video_paths):
    """Download multiple videos to local machine"""

```



```

print("Downloading videos to local machine...")

for video_path in video_paths:
    if os.path.exists(video_path):
        print(f"Downloading: {os.path.basename(video_path)}")
        files.download(video_path)
    else:
        print(f"x File not found: {video_path}")

print("✓ All available videos downloaded!")

def display_performance_metrics_enhanced(frame_metrics, video_info):
    """Enhanced performance metrics with proper FPS labeling"""
    if not frame_metrics:
        print("No metrics available")
        return

    ssim_scores = [m['ssim'] for m in frame_metrics]
    processing_times = [m['processing_time'] for m in frame_metrics]

    avg_ssim = np.mean(ssim_scores)
    avg_processing_time = np.mean(processing_times)

    # Calculate FPS metrics with proper labeling
    fps_during_processing = 1.0 / avg_processing_time if avg_processing_time > 0 else 0
    fps_original_video = video_info['fps']

    # Get output video FPS
    fps_enhanced_video = fps_original_video # Default fallback
    if os.path.exists('output_videos/enhanced_video.mp4'):
        cap = cv2.VideoCapture('output_videos/enhanced_video.mp4')
        if cap.isOpened():
            fps_enhanced_video = cap.get(cv2.CAP_PROP_FPS)
            cap.release()

    # Create comprehensive visualization
    fig, axes = plt.subplots(2, 3, figsize=(18, 12))

    # SSIM over time
    axes[0, 0].plot(ssim_scores, color='blue', linewidth=2)
    axes[0, 0].axhline(y=0.9, color='red', linestyle='--', label='Target (90%)')
    axes[0, 0].set_title(f'SSIM Over Time (Avg: {avg_ssim:.3f})')
    axes[0, 0].set_xlabel('Frame Number')
    axes[0, 0].set_ylabel('SSIM Score')
    axes[0, 0].legend()
    axes[0, 0].grid(True)

    # Enhanced FPS Comparison with proper labeling
    fps_categories = ['FPS During\nProcessing', 'Enhanced Video\nFPS', 'Original Video\nFPS']
    fps_values = [fps_during_processing, fps_enhanced_video, fps_original_video]
    fps_colors = ['orange', 'green', 'blue']

    bars = axes[0, 1].bar(fps_categories, fps_values, color=fps_colors, alpha=0.7)
    axes[0, 1].set_title('FPS Performance Comparison')
    axes[0, 1].set_ylabel('Frames Per Second (FPS)')
    axes[0, 1].grid(True, alpha=0.3)

    # Add value labels on bars
    for bar, value in zip(bars, fps_values):
        axes[0, 1].text(bar.get_x() + bar.get_width()/2, bar.get_height() + 1,
            f'{value:.1f}', ha='center', va='bottom', fontweight='bold')

    # Add target lines
    axes[0, 1].axhline(y=30, color='red', linestyle='--', alpha=0.7, label='30 FPS Target')
    axes[0, 1].axhline(y=60, color='purple', linestyle='--', alpha=0.7, label='60 FPS Target')
    axes[0, 1].legend()

    # SSIM histogram
    axes[0, 2].hist(ssim_scores, bins=20, alpha=0.7, color='blue', edgecolor='black')

```

```

axes[0, 2].axvline(x=0.9, color='red', linestyle='--', label='Target (90%)')
axes[0, 2].set_title('SSIM Distribution')
axes[0, 2].set_xlabel('SSIM Score')
axes[0, 2].set_ylabel('Frequency')
axes[0, 2].legend()
axes[0, 2].grid(True)

# Processing time histogram
axes[1, 0].hist(processing_times, bins=20, alpha=0.7, color='green', edgecolor='black')
axes[1, 0].set_title('Processing Time Distribution')
axes[1, 0].set_xlabel('Processing Time (seconds)')
axes[1, 0].set_ylabel('Frequency')
axes[1, 0].grid(True)

# Real-time capability pie chart
real_time_capable = fps_during_processing >= fps_original_video
axes[1, 1].pie([1 if real_time_capable else 0, 0 if real_time_capable else 1],
               labels=['Real-time\nCapable', 'Not Real-time'],
               colors=['green', 'red'],
               autopct='%1.0f%%',
               startangle=90)
axes[1, 1].set_title('Real-time Processing Capability')

# Enhanced Performance Summary
axes[1, 2].axis('off')
summary_text = f"""
ENHANCED PERFORMANCE SUMMARY

📊 Quality Metrics:
• Average SSIM: {avg_ssim:.4f} ({avg_ssim*100:.1f}%)
• SSIM Target (≥90%): {'✅ PASSED' if avg_ssim >= 0.9 else '❌ FAILED'}

⚡ Speed Metrics:
• FPS During Processing: {fps_during_processing:.1f}
• Enhanced Video FPS: {fps_enhanced_video:.1f}
• Original Video FPS: {fps_original_video:.1f}
• Real-time Capable: {'✅ YES' if real_time_capable else '❌ NO'}

🎯 Target Achievement:
• 30 FPS Processing: {'✅ PASSED' if fps_during_processing >= 30 else '❌ FAILED'}
• 60 FPS Processing: {'✅ PASSED' if fps_during_processing >= 60 else '❌ FAILED'}

📈 Frame Statistics:
• Total Frames: {len(frame_metrics)}
• Best SSIM: {max(ssim_scores):.4f}
• Worst SSIM: {min(ssim_scores):.4f}
"""

axes[1, 2].text(0.05, 0.95, summary_text, transform=axes[1, 2].transAxes,
               fontsize=10, verticalalignment='top', fontfamily='monospace',
               bbox=dict(boxstyle="round,pad=0.5", facecolor="lightblue", alpha=0.8))

plt.tight_layout()
plt.show()

return {
    'avg_ssim': avg_ssim,
    'fps_during_processing': fps_during_processing,
    'fps_enhanced_video': fps_enhanced_video,
    'fps_original_video': fps_original_video,
    'target_ssim_passed': avg_ssim >= 0.9,
    'target_fps_30_passed': fps_during_processing >= 30,
    'real_time_capable': real_time_capable
}

def complete_video_processing_pipeline():
    """Complete video processing pipeline with enhanced FPS reporting"""

    print("="*60)

```

```

print("PHASE 5: COMPLETE VIDEO PROCESSING PIPELINE")
print("="*60)

# Step 1: Upload video
print("\n📁 STEP 1: Upload your video file")
video_path = upload_video_file()

if not video_path:
    print("❌ Workflow stopped - no video uploaded")
    return

# Get video info
video_info = extract_video_info(video_path)
print(f"✓ Original Video FPS: {video_info['fps']}")
print(f"✓ Video Resolution: {video_info['width']}x{video_info['height']}")
print(f"✓ Total Frames: {video_info['frame_count']}")

# Step 2: Process frames
print("\n🔄 STEP 2: Processing video frames")
max_frames = min(100, video_info['frame_count'])
print(f"Processing first {max_frames} frames...")

original_frames, blurred_frames, enhanced_frames, frame_metrics = process_video_frames(
    video_path, student_model, max_frames=max_frames
)

if not original_frames:
    print("❌ Workflow stopped - frame processing failed")
    return

# Step 3: Create videos
print("\n🎬 STEP 3: Creating output videos")
output_dir = 'output_videos'
if not os.path.exists(output_dir):
    os.makedirs(output_dir)

video_paths = create_output_videos(original_frames, blurred_frames, enhanced_frames, video_info, output_dir)
comparison_path = create_side_by_side_comparison_video(
    original_frames, blurred_frames, enhanced_frames, video_info,
    os.path.join(output_dir, 'comparison_video.mp4')
)
video_paths.append(comparison_path)

# Step 4: Display videos
print("\n📺 STEP 4: Displaying videos in Colab")
video_titles = {
    'enhanced_video.mp4': 'Enhanced Video (Student Model)',
    'blurred_video.mp4': 'Blurred Video (Simulated Poor Network)',
    'comparison_video.mp4': 'Side-by-Side Comparison'
}

for video_path in video_paths:
    if os.path.exists(video_path):
        filename = os.path.basename(video_path)
        title = video_titles.get(filename, filename.replace('.mp4', '').replace('_', ' ').title())
        display_video_in_colab(video_path, title=title, width=800, height=400)

# Step 5: Enhanced performance analysis
print("\n📊 STEP 5: Enhanced Performance Analysis")
performance_summary = display_performance_metrics_enhanced(frame_metrics, video_info)

# Step 6: Download videos
print("\n💾 STEP 6: Downloading videos")
download_videos_to_local(video_paths)

print("\n" + "="*60)
print("🎉 PHASE 5 COMPLETE!")
print("="*60)
print(f"✅ Processed {len(original_frames)} frames")

```

```

print(f"✅ Original Video FPS: {performance_summary['fps_original_video']:.1f}")
print(f"✅ FPS During Processing: {performance_summary['fps_during_processing']:.1f}")
print(f"✅ Enhanced Video FPS: {performance_summary['fps_enhanced_video']:.1f}")
print(f"✅ Average SSIM: {performance_summary['avg_ssim']:.4f}")
print(f"✅ Real-time Capable: {'YES' if performance_summary['real_time_capable'] else 'NO'})")

return video_paths, performance_summary

print("✅ All Phase 5 functions defined successfully!")
print("Ready to run the complete video processing pipeline!")

```

➡️ ✅ All Phase 5 functions defined successfully!  
Ready to run the complete video processing pipeline!

Execute Phase 5

```

# Execute Phase 5: Complete Video Processing Pipeline
video_paths, performance = complete_video_processing_pipeline()

```

**Session Crashed after using all available RAM**

**Loading Your Saved Model in a New Colab Session**

Phase 1: Setup and Dependencies

```

# Install all required dependencies
!pip install torch torchvision torchaudio
!pip install opencv-python-headless
!pip install pillow
!pip install scikit-image
!pip install tqdm
!pip install ipywidgets

# Fix for PyTorch 2.6+ compatibility
!pip install huggingface_hub==0.25.2

# Import all required libraries
import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
from torch.utils.data import Dataset, DataLoader
import torchvision.transforms as transforms
import cv2
import numpy as np
from PIL import Image
import os
import zipfile
import urllib.request
from tqdm import tqdm
import matplotlib.pyplot as plt
from skimage.metrics import structural_similarity as ssim
import json
import time
from google.colab import files
import shutil
import glob
import base64
from IPython.display import display, HTML
import ipywidgets as widgets

# Set device
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

```

```
print(f"Using device: {device}")
print(f"CUDA available: {torch.cuda.is_available()}")
if torch.cuda.is_available():
    print(f"GPU: {torch.cuda.get_device_name(0)}")
```

```
Requirement already satisfied: torch in /usr/local/lib/python3.11/dist-packages (2.6.0+cu124)
Requirement already satisfied: torchvision in /usr/local/lib/python3.11/dist-packages (0.21.0+cu124)
Requirement already satisfied: torchaudio in /usr/local/lib/python3.11/dist-packages (2.6.0+cu124)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from torch) (3.18.0)
Requirement already satisfied: typing-extensions>=4.10.0 in /usr/local/lib/python3.11/dist-packages (from torch) (4.12.0)
Requirement already satisfied: networkx in /usr/local/lib/python3.11/dist-packages (from torch) (3.5)
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.11/dist-packages (from torch) (3.1.6)
Requirement already satisfied: fsspec in /usr/local/lib/python3.11/dist-packages (from torch) (2025.3.2)
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch) (12.4.127)
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch) (12.4.127)
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch) (12.4.127)
Requirement already satisfied: nvidia-cudnn-cu12==9.1.0.70 in /usr/local/lib/python3.11/dist-packages (from torch) (9.1.0.70)
Requirement already satisfied: nvidia-cublas-cu12==12.4.5.8 in /usr/local/lib/python3.11/dist-packages (from torch) (12.4.5.8)
Requirement already satisfied: nvidia-cufft-cu12==11.2.1.3 in /usr/local/lib/python3.11/dist-packages (from torch) (11.2.1.3)
Requirement already satisfied: nvidia-curand-cu12==10.3.5.147 in /usr/local/lib/python3.11/dist-packages (from torch) (10.3.5.147)
Requirement already satisfied: nvidia-cusolver-cu12==11.6.1.9 in /usr/local/lib/python3.11/dist-packages (from torch) (11.6.1.9)
Requirement already satisfied: nvidia-cusparselt-cu12==0.6.2 in /usr/local/lib/python3.11/dist-packages (from torch) (0.6.2)
Requirement already satisfied: nvidia-cusparse-cu12==12.3.1.170 in /usr/local/lib/python3.11/dist-packages (from torch) (12.3.1.170)
Requirement already satisfied: nvidia-nccl-cu12==2.21.5 in /usr/local/lib/python3.11/dist-packages (from torch) (2.21.5)
Requirement already satisfied: nvidia-nvtx-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch) (12.4.127)
Requirement already satisfied: nvidia-nvjitlink-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch) (12.4.127)
Requirement already satisfied: triton==3.2.0 in /usr/local/lib/python3.11/dist-packages (from torch) (3.2.0)
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.11/dist-packages (from torch) (1.13.1)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from sympy==1.13.1) (1.37.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from torchvision) (2.0.2)
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in /usr/local/lib/python3.11/dist-packages (from torchvision) (11.0.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.11/dist-packages (from Jinja2->torch) (3.0.2)
Requirement already satisfied: opencv-python-headless in /usr/local/lib/python3.11/dist-packages (4.11.0.86)
Requirement already satisfied: numpy>=1.21.2 in /usr/local/lib/python3.11/dist-packages (from opencv-python-headless) (2.0.2)
Requirement already satisfied: pillow in /usr/local/lib/python3.11/dist-packages (11.2.1)
Requirement already satisfied: scikit-image in /usr/local/lib/python3.11/dist-packages (0.25.2)
Requirement already satisfied: numpy>=1.24 in /usr/local/lib/python3.11/dist-packages (from scikit-image) (2.0.2)
Requirement already satisfied: scipy>=1.11.4 in /usr/local/lib/python3.11/dist-packages (from scikit-image) (1.14.1)
Requirement already satisfied: networkx>=3.0 in /usr/local/lib/python3.11/dist-packages (from scikit-image) (3.5)
Requirement already satisfied: pillow>=10.1 in /usr/local/lib/python3.11/dist-packages (from scikit-image) (11.0.0)
Requirement already satisfied: imageio!=2.35.0,>=2.33 in /usr/local/lib/python3.11/dist-packages (from scikit-image) (2.36.0)
Requirement already satisfied: tifffile>=2022.8.12 in /usr/local/lib/python3.11/dist-packages (from scikit-image) (2024.12.2)
Requirement already satisfied: packaging>=21 in /usr/local/lib/python3.11/dist-packages (from scikit-image) (25.0)
Requirement already satisfied: lazy-loader>=0.4 in /usr/local/lib/python3.11/dist-packages (from scikit-image) (0.4.0)
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (4.67.1)
Requirement already satisfied: ipywidgets in /usr/local/lib/python3.11/dist-packages (7.7.1)
Requirement already satisfied: ipykernel>=4.5.1 in /usr/local/lib/python3.11/dist-packages (from ipywidgets) (6.59.0)
Requirement already satisfied: ipython-genutils<0.2.0 in /usr/local/lib/python3.11/dist-packages (from ipywidgets) (0.1.0)
Requirement already satisfied: traitlets>=4.3.1 in /usr/local/lib/python3.11/dist-packages (from ipywidgets) (4.3.1)
Requirement already satisfied: widgetsnbextension<=3.6.0 in /usr/local/lib/python3.11/dist-packages (from ipywidgets) (3.6.0)
Requirement already satisfied: ipython>=4.0.0 in /usr/local/lib/python3.11/dist-packages (from ipywidgets) (7.34.0)
Requirement already satisfied: jupyterlab-widgets>=1.0.0 in /usr/local/lib/python3.11/dist-packages (from ipykernel>=4.5.1->ipywidgets) (1.0.0)
Requirement already satisfied: debugpy>=1.0 in /usr/local/lib/python3.11/dist-packages (from ipykernel>=4.5.1->ipywidgets) (1.8.1)
Requirement already satisfied: jupyter-client>=6.1.12 in /usr/local/lib/python3.11/dist-packages (from ipykernel>=4.5.1->ipywidgets) (6.4.0)
Requirement already satisfied: matplotlib-inline>=0.1 in /usr/local/lib/python3.11/dist-packages (from ipykernel>=4.5.1->ipywidgets) (0.1.7)
Requirement already satisfied: nest-asyncio in /usr/local/lib/python3.11/dist-packages (from ipykernel>=4.5.1->ipywidgets) (1.6.0)
Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from ipykernel>=4.5.1->ipywidgets) (25.0)
Requirement already satisfied: psutil in /usr/local/lib/python3.11/dist-packages (from ipykernel>=4.5.1->ipywidgets) (6.1.0)
Requirement already satisfied: pyzmq>=17 in /usr/local/lib/python3.11/dist-packages (from ipykernel>=4.5.1->ipywidgets) (26.2.1)
Requirement already satisfied: tornado>=6.1 in /usr/local/lib/python3.11/dist-packages (from ipykernel>=4.5.1->ipywidgets) (6.4.2)
Requirement already satisfied: setuptools>=18.5 in /usr/local/lib/python3.11/dist-packages (from ipython>=4.0->ipykernel>=4.5.1->ipywidgets) (75.8.2)
Requirement already satisfied: jedi>=0.16 in /usr/local/lib/python3.11/dist-packages (from ipython>=4.0->ipykernel>=4.5.1->ipywidgets) (0.19.2)
Requirement already satisfied: decorator in /usr/local/lib/python3.11/dist-packages (from ipython>=4.0->ipykernel>=4.5.1->ipywidgets) (5.1.1)
```

## Phase 2: Load Pre-trained Student Model

```
# Student Model Architecture (same as training)
class StudentModel(nn.Module):
    def __init__(self, scale_factor=4):
        super(StudentModel, self).__init__()
        self.scale_factor = scale_factor
```

```

# Feature extraction layers
self.conv1 = nn.Conv2d(3, 32, kernel_size=3, padding=1)
self.conv2 = nn.Conv2d(32, 64, kernel_size=3, padding=1)
self.conv3 = nn.Conv2d(64, 64, kernel_size=3, padding=1)

# Residual blocks
self.res_blocks = nn.ModuleList([
    self._make_residual_block(64) for _ in range(4)
])

# Upsampling layers
self.upsample = nn.Sequential(
    nn.Conv2d(64, 64 * (scale_factor ** 2), kernel_size=3, padding=1),
    nn.PixelShuffle(scale_factor),
    nn.Conv2d(64, 3, kernel_size=3, padding=1)
)

self.relu = nn.ReLU(inplace=True)
self.tanh = nn.Tanh()

def _make_residual_block(self, channels):
    return nn.Sequential(
        nn.Conv2d(channels, channels, kernel_size=3, padding=1),
        nn.ReLU(inplace=True),
        nn.Conv2d(channels, channels, kernel_size=3, padding=1)
    )

def forward(self, x):
    # Feature extraction
    x1 = self.relu(self.conv1(x))
    x2 = self.relu(self.conv2(x1))
    x3 = self.relu(self.conv3(x2))

    # Residual blocks
    residual = x3
    for res_block in self.res_blocks:
        out = res_block(residual)
        residual = residual + out

    # Upsampling
    out = self.upsample(residual)
    out = self.tanh(out)

    return out

# Initialize student model
student_model = StudentModel(scale_factor=4).to(device)

# Load your pre-trained model
def load_pretrained_student_model():
    """Load the pre-trained student model from checkpoint"""

    # Option 1: If model file exists in current directory
    checkpoint_path = 'best_student_model.pth'

    if os.path.exists(checkpoint_path):
        try:
            checkpoint = torch.load(checkpoint_path, map_location=device, weights_only=False)
            student_model.load_state_dict(checkpoint['model_state_dict'])
            student_model.eval()
            print(f"✅ Loaded pre-trained model from {checkpoint_path}")
            print(f"    Training epoch: {checkpoint.get('epoch', 'unknown')}")
            print(f"    SSIM score: {checkpoint.get('ssim', 'unknown'):.4f}")
            return True
        except Exception as e:
            print(f"❌ Error loading model: {e}")
            return False
    else:
        # Option 2: Upload model file

```

```

print("Model file not found. Please upload your trained model:")
uploaded = files.upload()

if uploaded:
    model_file = list(uploaded.keys())[0]
    try:
        checkpoint = torch.load(model_file, map_location=device, weights_only=False)
        student_model.load_state_dict(checkpoint['model_state_dict'])
        student_model.eval()
        print(f"✅ Loaded pre-trained model from {model_file}")
        print(f"    Training epoch: {checkpoint.get('epoch', 'unknown')}")
        print(f"    SSIM score: {checkpoint.get('ssim', 'unknown'):.4f}")
        return True
    except Exception as e:
        print(f"❌ Error loading uploaded model: {e}")
        return False
else:
    print("❌ No model file uploaded")
    return False

# Load the pre-trained model
model_loaded = load_pretrained_student_model()

if model_loaded:
    print(f"✅ Student model parameters: {sum(p.numel() for p in student_model.parameters()):,}")
    print("✅ Pre-trained model ready for video processing!")
else:
    print("❌ Failed to load pre-trained model. Please check your model file.")

```

📁 Model file not found. Please upload your trained model:

best\_stude...5epochs.pth

- **best\_student\_model\_25epochs.pth**(n/a) - 11366482 bytes, last modified: 7/9/2025 - 100% done
- Saving best\_student\_model\_25epochs.pth to best\_student\_model\_25epochs (3).pth
- ✅ Loaded pre-trained model from best\_student\_model\_25epochs (3).pth
  - Training epoch: 21
  - SSIM score: 0.7577
  - ✅ Student model parameters: 944,323
  - ✅ Pre-trained model ready for video processing!

### Phase 3: Complete Video Processing Pipeline

# Video processing functions with all requested features

```

def upload_video_file():
    """Upload video file from local machine"""
    print("Please upload a video file...")
    print("Supported formats: .mp4, .avi, .mov, .mkv")

    uploaded = files.upload()

    if not uploaded:
        print("No file uploaded!")
        return None

    video_path = list(uploaded.keys())[0]
    print(f"✓ Video uploaded: {video_path}")
    return video_path

def extract_video_info(video_path):
    """Extract basic video information"""
    cap = cv2.VideoCapture(video_path)
    if not cap.isOpened():
        return None

    fps = int(cap.get(cv2.CAP_PROP_FPS))
    width = int(cap.get(cv2.CAP_PROP_FRAME_WIDTH))
    height = int(cap.get(cv2.CAP_PROP_FRAME_HEIGHT))
    frame_count = int(cap.get(cv2.CAP_PROP_FRAME_COUNT))

```

```

duration = frame_count / fps if fps > 0 else 0

cap.release()

return {
    'fps': fps,
    'width': width,
    'height': height,
    'frame_count': frame_count,
    'duration': duration
}

def simulate_network_blur(frame, blur_factor=4):
    """Simulate poor network conditions by downscaling and upscaling"""
    h, w = frame.shape[:2]
    small_frame = cv2.resize(frame, (w//blur_factor, h//blur_factor), interpolation=cv2.INTER_CUBIC)
    blurred_frame = cv2.resize(small_frame, (w, h), interpolation=cv2.INTER_CUBIC)
    return blurred_frame

def enhance_frame_with_student(frame, student_model):
    """Enhance a single frame using the trained student model"""
    try:
        frame_rgb = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)
        frame_tensor = torch.from_numpy(frame_rgb).float() / 255.0
        frame_tensor = frame_tensor.permute(2, 0, 1).unsqueeze(0).to(device)

        student_model.eval()
        with torch.no_grad():
            enhanced_tensor = student_model(frame_tensor)

        enhanced_frame = enhanced_tensor.squeeze(0).permute(1, 2, 0).cpu().numpy()
        enhanced_frame = np.clip(enhanced_frame * 255, 0, 255).astype(np.uint8)
        enhanced_frame = cv2.cvtColor(enhanced_frame, cv2.COLOR_RGB2BGR)

        if enhanced_frame.shape[:2] != frame.shape[:2]:
            enhanced_frame = cv2.resize(enhanced_frame, (frame.shape[1], frame.shape[0]))

        return enhanced_frame

    except Exception as e:
        print(f"Enhancement failed: {e}")
        return frame

def process_video_frames(video_path, student_model, max_frames=None):
    """Process video frames: extract, blur, enhance"""
    print(f"Processing video: {video_path}")

    video_info = extract_video_info(video_path)
    if video_info is None:
        print("x Failed to read video file")
        return None, None, None, None

    cap = cv2.VideoCapture(video_path)
    original_frames = []
    blurred_frames = []
    enhanced_frames = []
    frame_metrics = []

    frame_count = 0
    max_process = max_frames if max_frames else video_info['frame_count']

    pbar = tqdm(total=min(max_process, video_info['frame_count']), desc="Processing frames")

    while True:
        ret, frame = cap.read()
        if not ret or frame_count >= max_process:
            break

        original_frames.append(frame.copy())

```